



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Experiment No. 3
Implementation of Bayes classifier for recommendation
Date of Performance:
Date of Submission:
Marks:
Sign:



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

Aim: Implementation of Bayes classifier for recommendation

Objective: Able to perform naïve bayes classifier for implementation recommendations system.

Theory: Naive Bayes classifiers are a collection of classification algorithms based on Bayes' Theorem. It is not a single algorithm but a family of algorithms where all of them share a common principle, i.e. every pair of features being classified is independent of each other. To start with, let us consider a dataset.

One of the most simple and effective classification algorithms, the Naïve Bayes classifier aids in the rapid development of machine learning models with rapid prediction capabilities.

Naïve Bayes algorithm is used for classification problems. It is highly used in text classification. In text classification tasks, data contains high dimension (as each word represent one feature in the data). It is used in spam filtering, sentiment detection, rating classification etc. The advantage of using naïve Bayes is its speed. It is fast and making prediction is easy with high dimension of data.

Implementation:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import os
for dirname, _, filenames in
os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

import warnings
warnings.filterwarnings('ignore')
data = '/kaggle/input/adult-dataset/adult.csv'
df = pd.read_csv(data, header=None, sep=',')
col_names = ['age', 'workclass', 'fnlwgt', 'education',
'education_num', 'marital_status', 'occupation',
'relationship',
'race', 'sex', 'capital_gain', 'capital_loss',
'hours_per_week', 'native_country', 'income']
df.columns = col_names
categorical = [var for var in df.columns if
df[var].dtype=='O']
```

```
print('There are {} categorical
variables\n'.format(len(categorical)))
print('The categorical variables are :\n\n',
categorical)
df[categorical].isnull().sum()
for var in categorical:
    print(df[var].value_counts())
for var in categorical:
    print(df[var].value_counts()/np.float(len(df)))
df.workclass.unique()
df.workclass.value_counts()
df['workclass'].replace('?', np.NaN, inplace=True)
df.workclass.value_counts()
df.occupation.unique()
df.occupation.value_counts()
df['occupation'].replace('?', np.NaN, inplace=True)
df.occupation.value_counts()
df.native_country.unique()
df.native_country.value_counts()
df['native_country'].replace('?', np.NaN,
inplace=True)
df.native_country.value_counts()
df[categorical].isnull().sum()
for var in categorical:
```



Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
print(var, ' contains ', len(df[var].unique()), '
labels')
numerical = [var for var in df.columns if
df[var].dtype!='O']
print('There are {} numerical
variables\n'.format(len(numerical)))
print('The numerical variables are :', numerical)
df[numerical].head()
df[numerical].isnull().sum()
X = df.drop(['income'], axis=1)
y = df['income']
from sklearn.model_selection import
train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,
y, test_size = 0.3, random_state = 0)
X_train.shape, X_test.shape
X_train.dtypes
categorical = [col for col in X_train.columns if
X_train[col].dtype == 'O']
numerical = [col for col in X_train.columns if
X_train[col].dtype != 'O']
X_train[categorical].isnull().mean()
for col in categorical:
    if X_train[col].isnull().mean()>0:
        print(col, (X_train[col].isnull().mean()))
for df2 in [X_train, X_test]:

df2['workclass'].fillna(X_train['workclass'].mode()[
0], inplace=True)

df2['occupation'].fillna(X_train['occupation'].mode(
))[0], inplace=True)

df2['native_country'].fillna(X_train['native_country
'].mode()[0], inplace=True)
X_train[categorical].isnull().sum()
X_test[categorical].isnull().sum()
X_train.isnull().sum()
X_test.isnull().sum()
import category_encoders as ce
encoder = ce.OneHotEncoder(cols=['workclass',
'education', 'marital_status', 'occupation',
'relationship',
'race', 'sex', 'native_country'])
X_train = encoder.fit_transform(X_train)
X_test = encoder.transform(X_test)
cols = X_train.columns
from sklearn.preprocessing import RobustScaler
scaler = RobustScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
X_train = pd.DataFrame(X_train, columns=[cols])
X_test = pd.DataFrame(X_test, columns=[cols])
from sklearn.naive_bayes import GaussianNB
gnb = GaussianNB()
gnb.fit(X_train, y_train)
```

```
y_pred = gnb.predict(X_test)
from sklearn.metrics import accuracy_score
print('Model accuracy score: {0:0.4f}'.
format(accuracy_score(y_test, y_pred)))
y_pred_train = gnb.predict(X_train)
print('Training-set accuracy score: {0:0.4f}'.
format(accuracy_score(y_train, y_pred_train)))
print('Training set score:
{:.4f}'.format(gnb.score(X_train, y_train)))
print('Test set score:
{:.4f}'.format(gnb.score(X_test, y_test)))
y_test.value_counts()
null_accuracy = (7407/(7407+2362))
print('Null accuracy score: {0:0.4f}'.
format(null_accuracy))
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)
print('Confusion matrix\n\n', cm)
print('\nTrue Positives(TP) = ', cm[0,0])
print('\nTrue Negatives(TN) = ', cm[1,1])
print('\nFalse Positives(FP) = ', cm[0,1])
print('\nFalse Negatives(FN) = ', cm[1,0])
cm_matrix = pd.DataFrame(data=cm,
columns=['Actual Positive:1', 'Actual Negative:0'],
index=['Predict Positive:1',
'Predict Negative:0'])
sns.heatmap(cm_matrix, annot=True, fmt='d',
cmap='YlGnBu')
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
TP = cm[0,0]
TN = cm[1,1]
FP = cm[0,1]
FN = cm[1,0]
classification_accuracy = (TP + TN) / float(TP +
TN + FP + FN)
print('Classification accuracy :
{0:0.4f}'.format(classification_accuracy))
classification_error = (FP + FN) / float(TP + TN +
FP + FN)
print('Classification error :
{0:0.4f}'.format(classification_error))
precision = TP / float(TP + FP)
print('Precision : {0:0.4f}'.format(precision))
recall = TP / float(TP + FN)
print('Recall or Sensitivity :
{0:0.4f}'.format(recall))
true_positive_rate = TP / float(TP + FN)
print('True Positive Rate :
{0:0.4f}'.format(true_positive_rate))
false_positive_rate = FP / float(FP + TN)
print('False Positive Rate :
{0:0.4f}'.format(false_positive_rate))
specificity = TN / (TN + FP)
print('Specificity : {0:0.4f}'.format(specificity))
y_pred_prob = gnb.predict_proba(X_test)[0:10]
```




Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

```
y_pred_prob_df =  
pd.DataFrame(data=y_pred_prob, columns=['Prob  
of - <=50K', 'Prob of - >50K'])  
gnb.predict_proba(X_test)[0:10, 1]  
y_pred1 = gnb.predict_proba(X_test)[: , 1]  
plt.rcParams['font.size'] = 12  
plt.hist(y_pred1, bins = 10)  
plt.title('Histogram of predicted probabilities of  
salaries >50K')  
plt.xlim(0,1)  
plt.xlabel('Predicted probabilities of salaries >50K')  
plt.ylabel('Frequency')  
from sklearn.metrics import roc_curve  
fpr, tpr, thresholds = roc_curve(y_test, y_pred1,  
pos_label = '>50K')  
plt.figure(figsize=(6,4))  
plt.plot(fpr, tpr, linewidth=2)  
plt.plot([0,1], [0,1], 'k--')  
plt.rcParams['font.size'] = 12  
plt.title('ROC curve for Gaussian Naive Bayes  
Classifier for Predicting Salaries')  
plt.xlabel('False Positive Rate (1 - Specificity)')  
plt.ylabel('True Positive Rate (Sensitivity)')  
plt.show()  
from sklearn.metrics import roc_auc_score  
ROC_AUC = roc_auc_score(y_test, y_pred1)  
print('ROC AUC : {:.4f}'.format(ROC_AUC))  
from sklearn.model_selection import  
cross_val_score  
Cross_validated_ROC_AUC =  
cross_val_score(gnb, X_train, y_train, cv=5,  
scoring='roc_auc').mean()  
print('Cross validated ROC AUC :  
{:.4f}'.format(Cross_validated_ROC_AUC))  
from sklearn.model_selection import  
cross_val_score  
scores = cross_val_score(gnb, X_train, y_train, cv  
= 10, scoring='accuracy')  
print('Cross-validation scores: {}'.format(scores))  
print('Average cross-validation score:  
{:.4f}'.format(scores.mean()))
```

Output:

```
Model accuracy score: 0.8083  
Training-set accuracy score: 0.8067  
Training set score: 0.8067  
Test set score: 0.8083  
Null accuracy score: 0.7582  
Confusion matrix  
[[5999 1408]  
 [ 465 1897]]  
True Positives (TP) = 5999  
True Negatives (TN) = 1897  
False Positives (FP) = 1408  
False Negatives (FN) = 465
```







Vidyavardhini's College of Engineering and Technology

Department of Artificial Intelligence & Data Science

	precision	recall	f1-score	support
<=50K	0.93	0.81	0.86	7407
>50K	0.57	0.80	0.67	2362
accuracy			0.81	9769
macro avg	0.75	0.81	0.77	9769
weighted avg	0.84	0.81	0.82	9769

Classification accuracy : 0.8083
Classification error : 0.1917
Precision : 0.8099
Recall or Sensitivity : 0.9281
True Positive Rate : 0.9281
False Positive Rate : 0.4260
Specificity : 0.5740



ROC AUC : 0.8941
Cross validated ROC AUC : 0.8938
Cross-validation scores:[0.81359649 0.80438596 0.81184211 0.80517771 0.79640193 0.79684072
0.81044318 0.81175954 0.80210619 0.81035996]
Average cross-validation score: 0.8063

Conclusion:

In conclusion, implementing the Bayes classifier for recommendation systems offers a probabilistic framework for personalized suggestions. Leveraging Bayesian inference, it efficiently predicts user preferences based on past behavior and item attributes. Its simplicity and effectiveness make it a valuable tool for enhancing recommendation accuracy and user satisfaction.